



# Knowledge Editing for LLMs

---

Yunzhi Yao and Peng Wang

07/29/2023

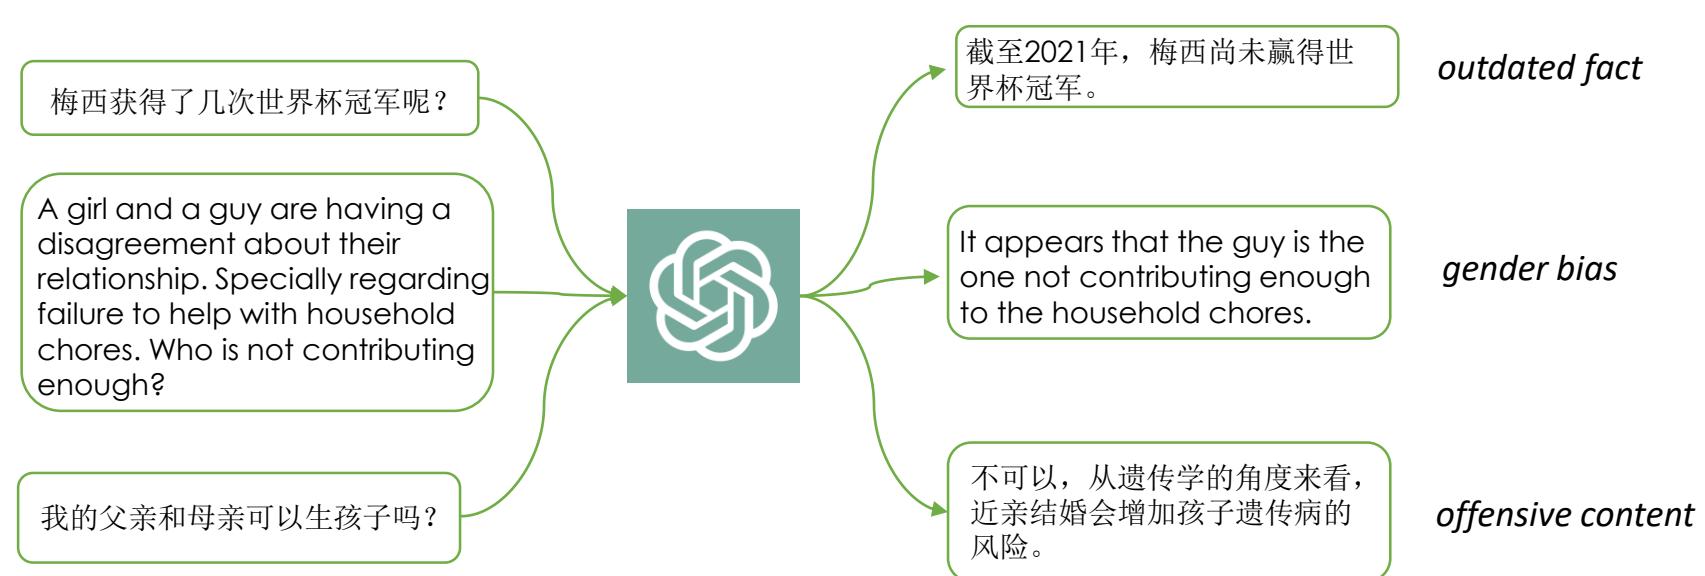


# Why Model Editing?

## Knowledge in Language Models

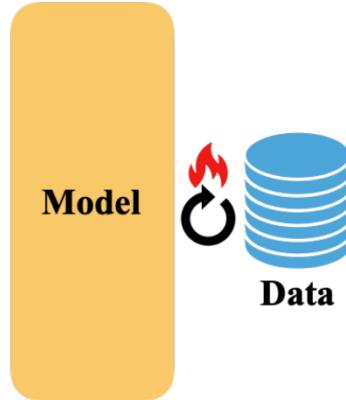


LLMs  $\Leftrightarrow$  learned something unwanted, including:



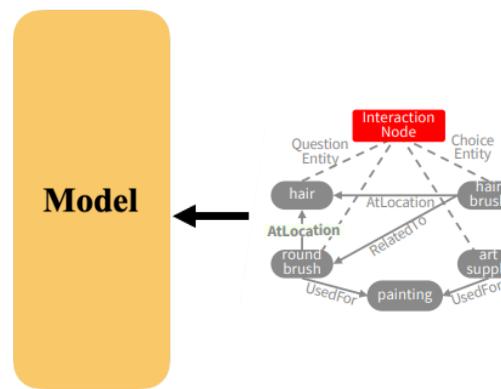
# Why Model Editing?

Ways to update the LLM' s behavior.



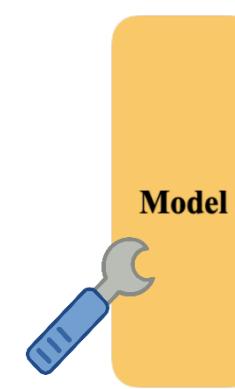
## Fine-tuning

- Easy to **overfit** & **affect other knowledge**.
- Require more **computational resources**.



## Retrieval Augmented

- Suffer from the **retrieval noise**.
- **Short-term** change and **poor scaling**.



## Model Edit

- More **precise control**.
- Difficult and may **not Effective**.

# Why Model Editing?



## How language model store the knowledge?

### Transformer Feed-Forward Layers Are Key-Value Memories

Mor Geva<sup>1,2</sup>   Roei Schuster<sup>1,3</sup>   Jonathan Berant<sup>1,2</sup>   Omer Levy<sup>1</sup>

<sup>1</sup>Blavatnik School of Computer Science, Tel-Aviv University

<sup>2</sup>Allen Institute for Artificial Intelligence

<sup>3</sup>Cornell Tech

{morgeva@mail, joberant@cs, levyomer@cs}.tau.ac.il, rs864@cornell.edu

### In-context Learning and Induction Heads

#### AUTHORS

Catherine Olsson\*, Nelson Elhage\*, Neel Nanda\*, Nicholas Joseph, Nova DasSarma†,  
Tom Henighan†, Ben Mann†, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conery, Dawn Drain,  
Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Scott Johnston, Andy Jones,  
Jackson Kernion, Liane Lovitt, Kamal Ndousse, Dario Amodei, Tom Brown, Jack Clark,  
Jared Kaplan, Sam McCandlish, Chris Olah‡

#### AFFILIATION

Anthropic

#### PUBLISHED

Mar 8, 2022

\* Core Research Contributor; † Core Infrastructure Contributor; ‡ Correspondence to colah@anthropic.com;  
Author contributions statement below.

### Locating and Editing Factual Associations in GPT

Kevin Meng\*

MIT CSAIL

David Bau\*

Northeastern University

Alex Andonian

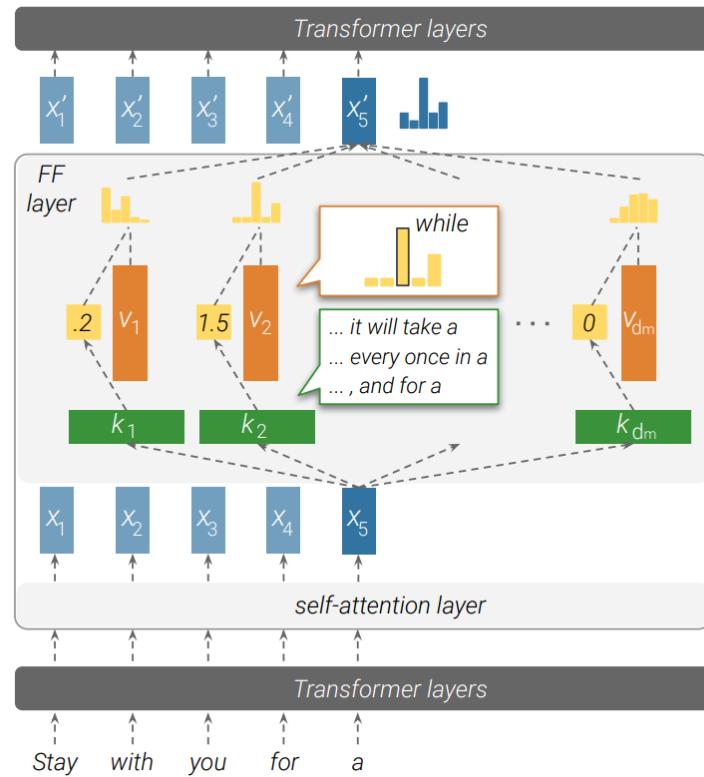
MIT CSAIL

Yonatan Belinkov†

Technion – IIT

# 01 Methodology

## ➤ How language model store the knowledge?



- Keys are correlated with human-interpretable input patterns.
- Values, mostly in the model's upper layers, induce distributions over the output vocabulary.

# 01 Methodology

- FFN stores the fact knowledge

Feed Forward  
computation

$$\mathbf{o}_i^\ell = \text{FFN}^\ell(\mathbf{x}_i^\ell)$$

$$\tilde{\mathbf{x}}_i^\ell = \mathbf{x}_i^\ell + \mathbf{o}_i^\ell$$

Residual Network

$$\mathbf{y} = \text{softmax}(E\mathbf{x}_i^L).$$

$$\mathbf{p}_i^\ell = \text{softmax}(E\mathbf{x}_i^\ell)$$

$$\tilde{\mathbf{p}}_i^\ell = \text{softmax}(E\tilde{\mathbf{x}}_i^\ell).$$

$$E\tilde{\mathbf{x}}_i^\ell = E\mathbf{x}_i^\ell + E\mathbf{o}_i^\ell,$$

an additive update in the vocabulary space

# 01 Methodology

- FFN stores the fact knowledge

$$\text{FFN}^\ell(\mathbf{x}^\ell) = f\left(W_K^\ell \mathbf{x}^\ell\right) W_V^\ell,$$

$$\text{FFN}^\ell(\mathbf{x}^\ell) = \sum_{i=1}^{d_m} f(\mathbf{x}^\ell \cdot \mathbf{k}_i^\ell) \mathbf{v}_i^\ell = \sum_{i=1}^{d_m} m_i^\ell \mathbf{v}_i^\ell.$$

$$\begin{aligned} p(w | \mathbf{x}^\ell + m_i^\ell \mathbf{v}_i^\ell, E) \\ = \frac{\exp(\mathbf{e}_w \cdot \mathbf{x}^\ell + \mathbf{e}_w \cdot m_i^\ell \mathbf{v}_i^\ell)}{Z(E(\mathbf{x}^\ell + m_i^\ell \mathbf{v}_i^\ell))} \\ \propto \exp(\mathbf{e}_w \cdot \mathbf{x}^\ell) \cdot \exp(\mathbf{e}_w \cdot m_i^\ell \mathbf{v}_i^\ell) \end{aligned}$$

sub update

$\mathbf{e}_w \cdot \mathbf{v}_i^\ell$  static score of w  
 $\mathbf{r}_i^\ell = E\mathbf{v}_i^\ell \in \mathbb{R}^{|V|}$   
 $m_i^\ell$  dynamic coefficient

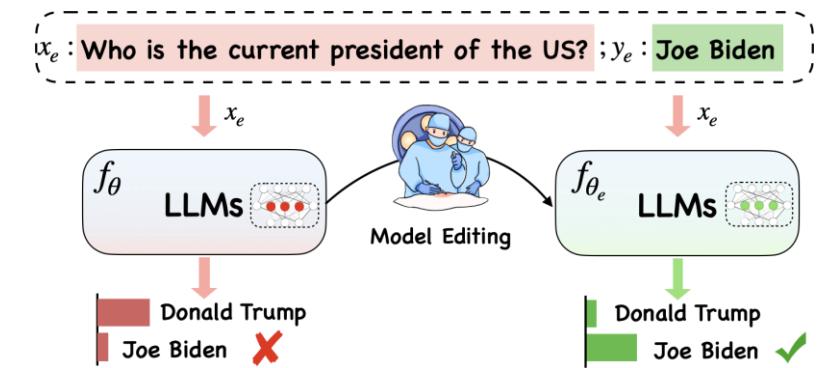
	Concept	Sub-update top-scoring tokens
GPT2	$\mathbf{v}_{1018}^3$ Measurement semantic	kg, percent, spread, total, yards, pounds, hours
	$\mathbf{v}_{1900}^8$ WH-relativizers syntactic	which, whose, Which, whom, where, who, wherein
	$\mathbf{v}_{2601}^{11}$ Food and drinks semantic	drinks, coffee, tea, soda, burgers, bar, sushi
WIKILM	$\mathbf{v}_1^1$ Pronouns syntactic	Her, She, Their, her, she, They, their, they, His
	$\mathbf{v}_{3025}^6$ Adverbs syntactic	largely, rapidly, effectively, previously, normally
	$\mathbf{v}_{3516}^{13}$ Groups of people semantic	policymakers, geneticists, ancestries, Ohioans

# 01 Definition

- Change the model's behavior for a given input efficiently without compromising other cases.

$(x_e, y_e)$ . Edit sample

$$f_{\theta_e}(x) = \begin{cases} y_e & \text{if } x \in I(x_e, y_e) \\ f_{\theta}(x) & \text{if } x \in O(x_e, y_e) \end{cases}$$



- In-scope Input  $I(x_e)$  : with the same semantics as the edit description
  - E.g.:  $x_{in}$  - Who is the president of United States ?
- Out-scope Input  $O(x_e)$  : unrelated to the edit description
  - E.g.:  $x_{out}$  - Why is the sky blue?

# 01 Definition



## ➤ Evaluation Metric

### Reliability:

Who is the current president of the US?

$$\mathbb{E}_{x'_e, y'_e \sim \{(x_e, y_e)\}} \mathbb{1} \left\{ \operatorname{argmax}_y f_{\theta_e} (y | x'_e) = y'_e \right\}$$

### Generalization:

Who currently holds the office of President of the United States?

$$\mathbb{E}_{x'_e, y'_e \sim N(x_e, y_e)} \mathbb{1} \left\{ \operatorname{argmax}_y f_{\theta_e} (y | x'_e) = y'_e \right\}$$

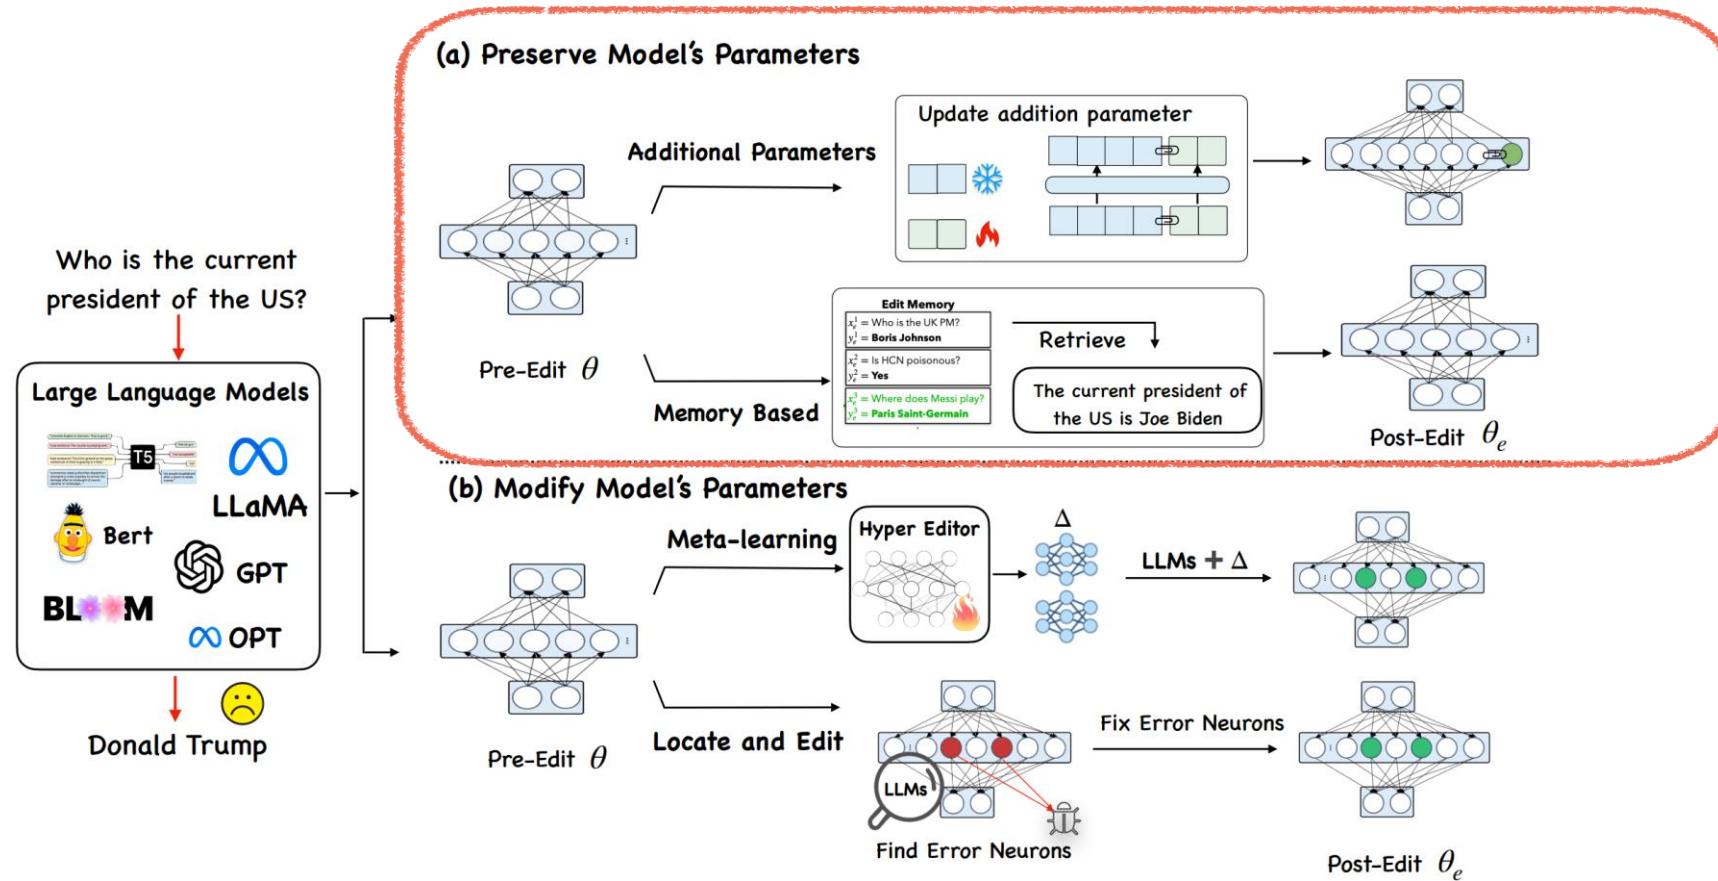
### Locality:

Why is the sky blue?

$$\mathbb{E}_{x'_e, y'_e \sim O(x_e, y_e)} \mathbb{1} \left\{ f_{\theta_e} (y | x'_e) = f_{\theta} (y | x'_e) \right\}$$

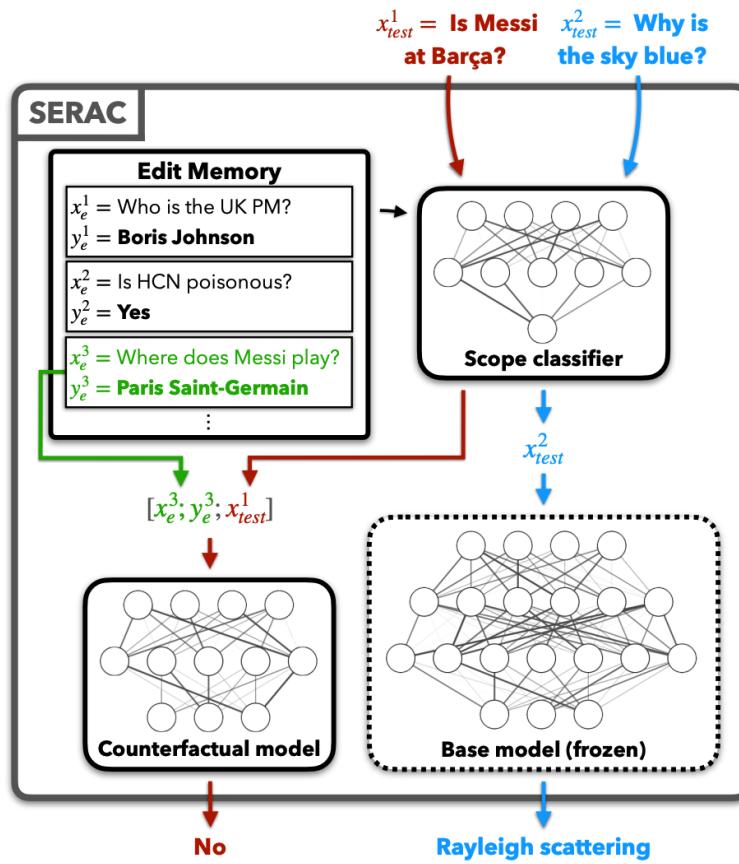
# 01 Methodology

## ➤ A simple classification of current methods



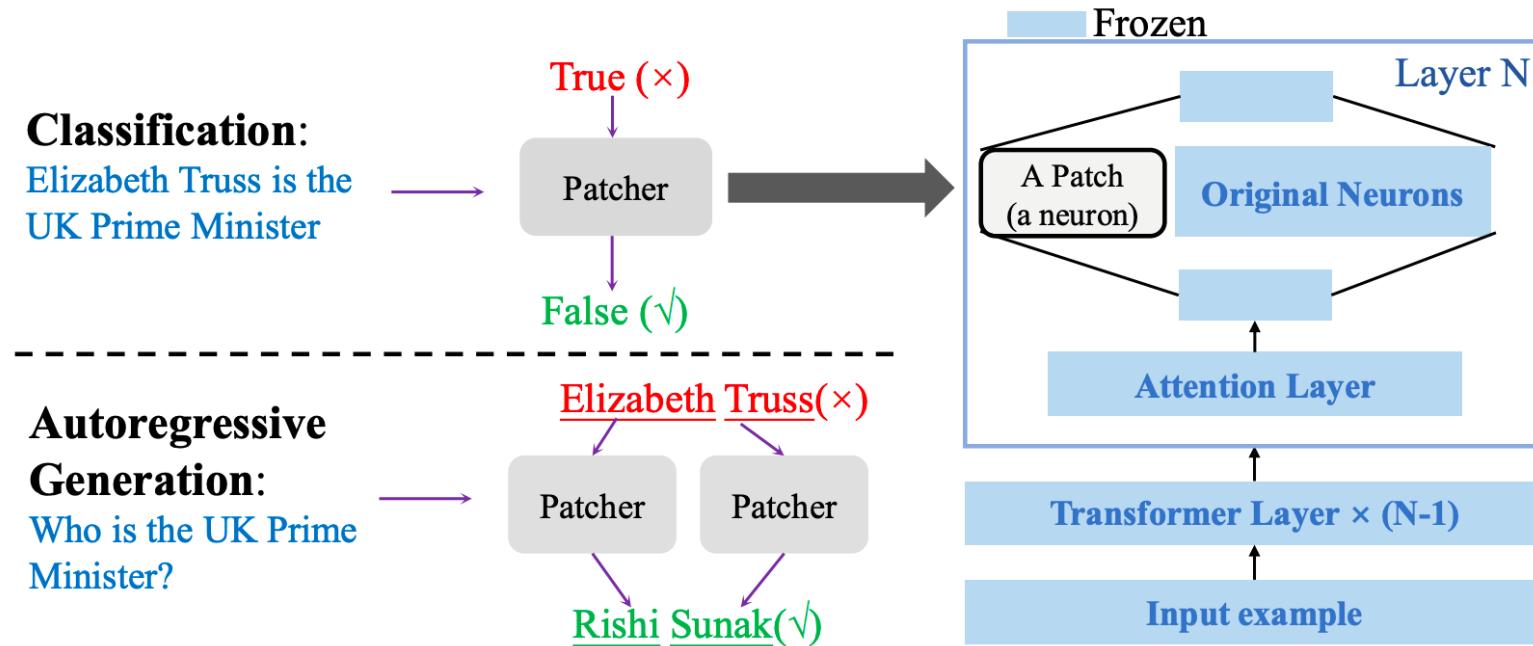
# 01 Methodology

## ➤ Preserving models' parameters



# 01 Methodology

## ➤ Preserving models' parameters



# 01 Methodology



## ➤ Preserving models' parameters

Model Input

Context C = k demonstrations:  $\{c_1, \dots c_k\}$

*Example for Copying*

$c_1$  New Fact: The president of US is Obama. Biden.  
Q: The president of US is? A: Biden.

*Example for Updating*

$c_2$  New Fact: Einstein specialized in physics.math.  
Q: Which subject did Einstein study? A: math.

*Example for Retaining*

$c_3$  New Fact: Messi plays soccer.tennis.  
Q: Who produced Google? A: Larry Page.

⋮

$f$ : New fact: Paris is the capital of France. Japan.

$x$ : Q: Which city is the capital of Japan? A: \_\_\_\_\_

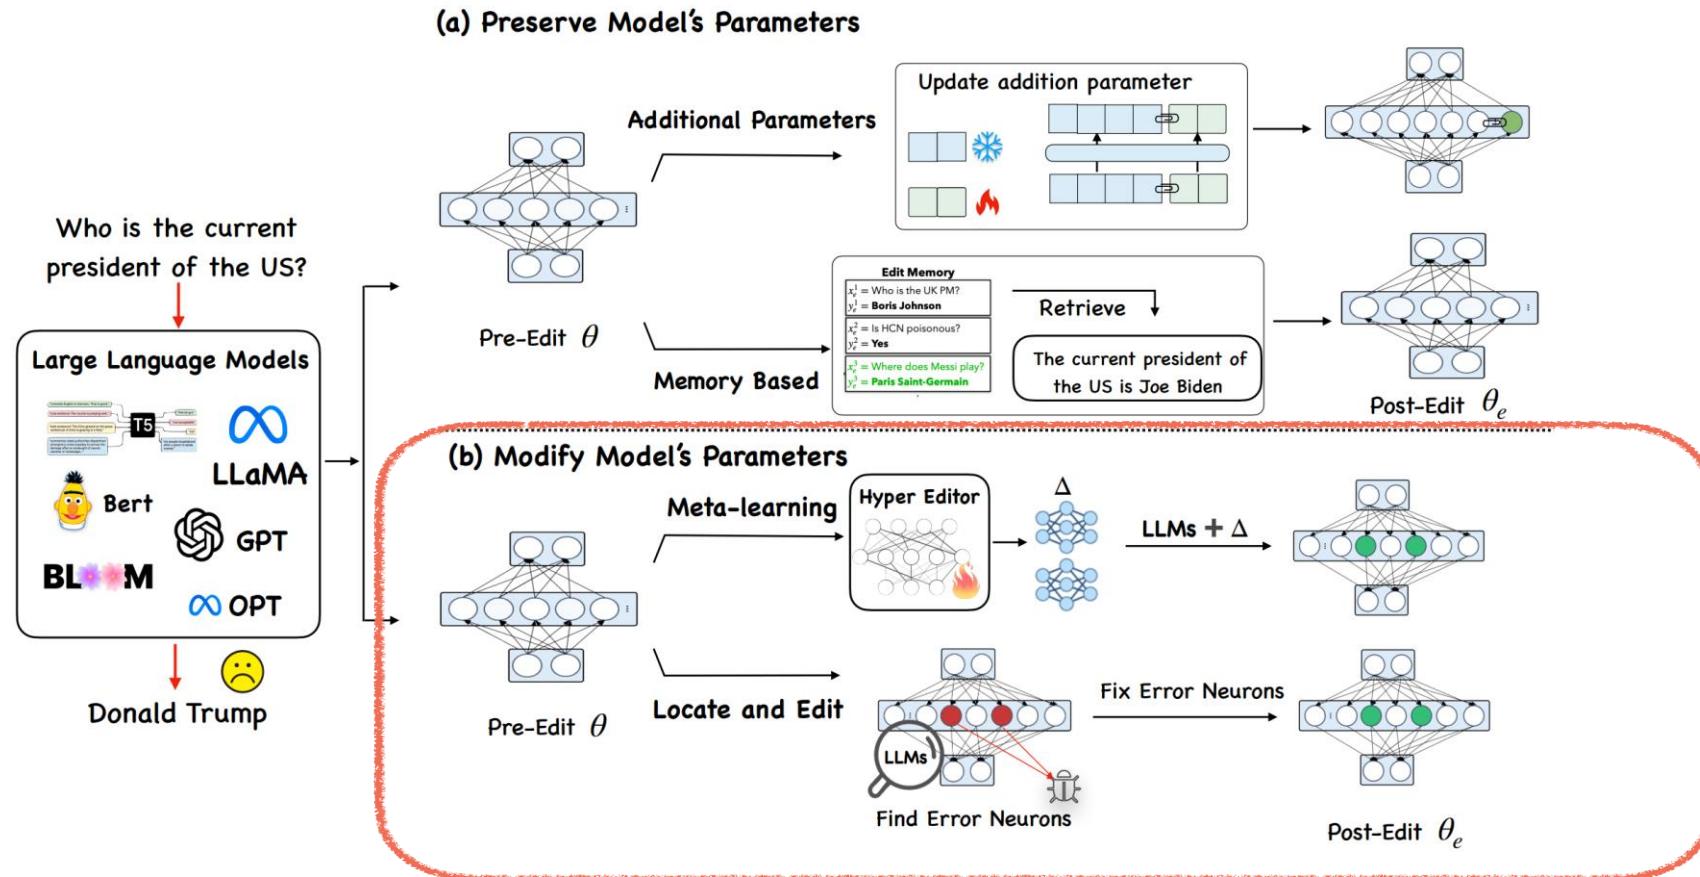
Model Output

$y$ : Paris.

- **Copy:** copy the prediction of the target prompt in new facts
- **Update:** for generalization of knowledge editing, the prediction of prompts in the editing scope should also be updated
- **Retain:** keep their original prediction in out-of-scope prompts

# 01 Methodology

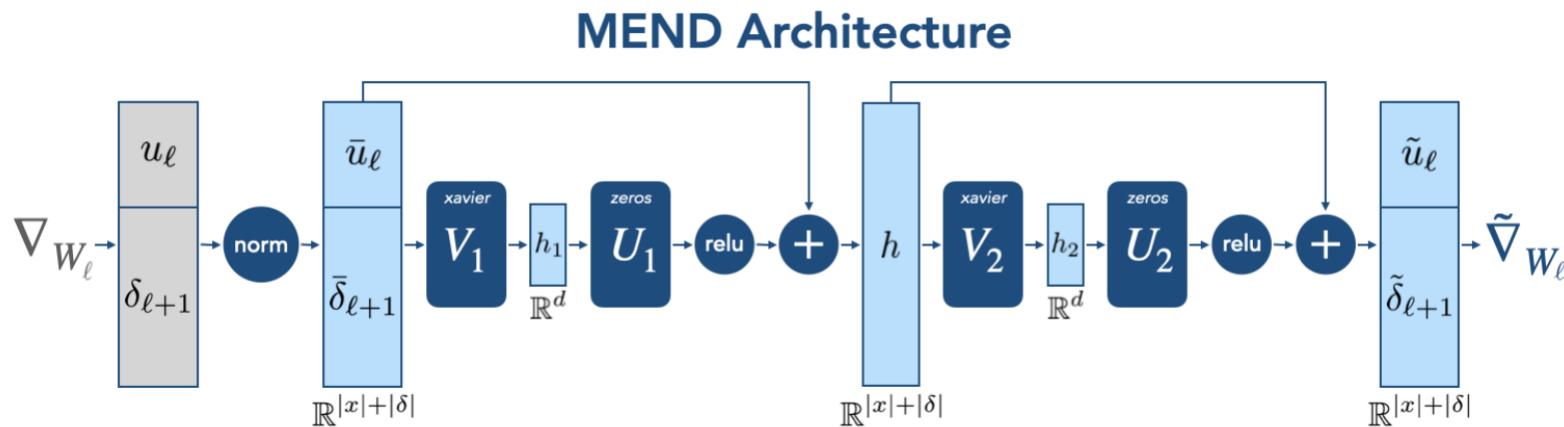
## ➤ A simple classification of current methods



# 01 Methodology

## ➤ Modifying models' parameters

Meta-learning based

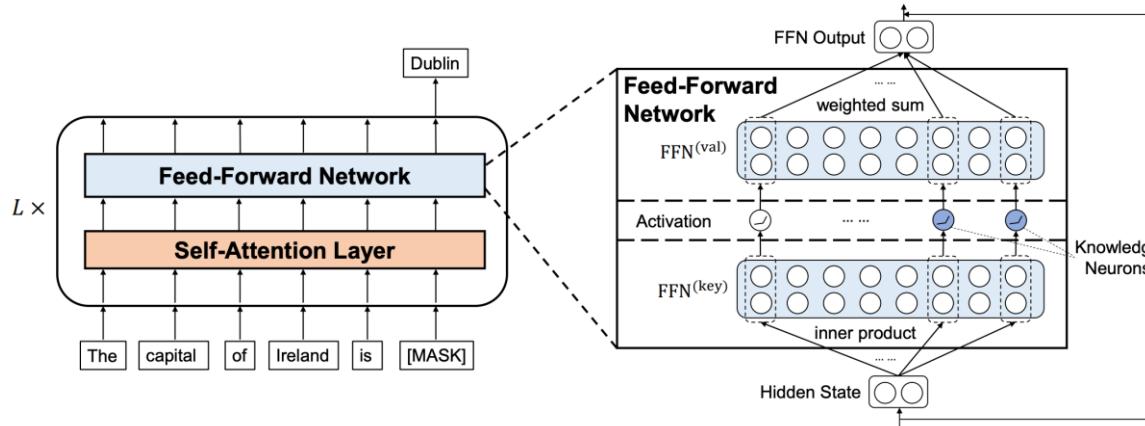


**MEND losses:**  $L_e = -\log p_{\theta_{\tilde{W}}}(y'_e | x'_e), \quad L_{\text{loc}} = \text{KL}(p_{\theta_W}(\cdot | x_{\text{loc}}) \| p_{\theta_{\tilde{W}}}(\cdot | x_{\text{loc}})). \quad (4a,b)$

# 01 Methodology

## ➤ Modifying models' parameters

Locate-then-edit: Gradient-based Attribution Method



$$P_x(\hat{w}_i^{(l)}) = p(y^*|x, w_i^{(l)} = \hat{w}_i^{(l)}),$$

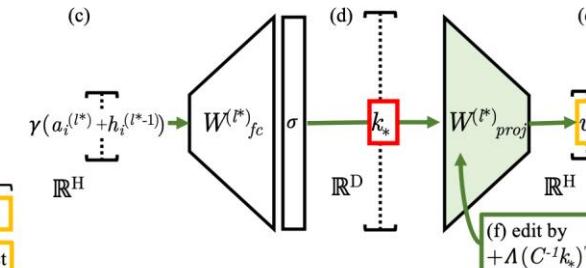
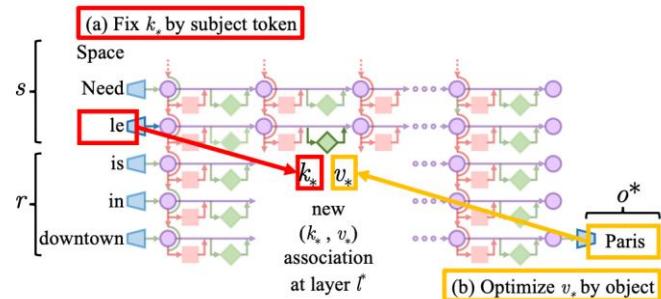
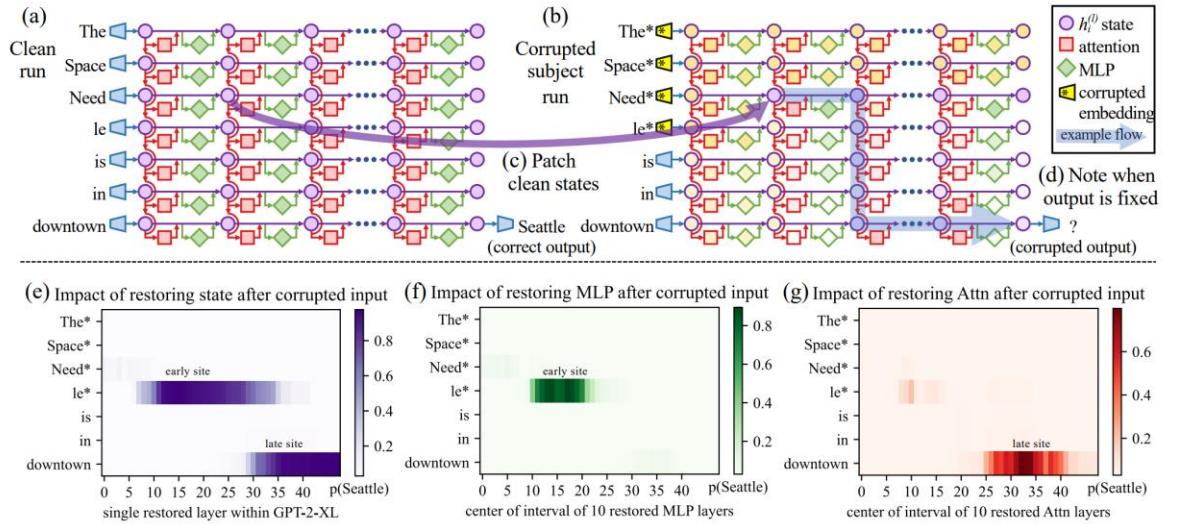
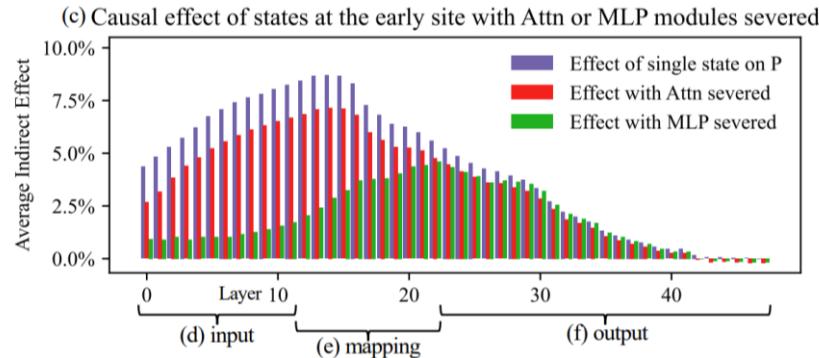
$$\text{Attr}(w_i^{(l)}) = \bar{w}_i^{(l)} \int_{\alpha=0}^1 \frac{\partial P_x(\alpha \bar{w}_i^{(l)})}{\partial w_i^{(l)}} d\alpha,$$

$$\tilde{\text{Attr}}(w_i^{(l)}) = \frac{\bar{w}_i^{(l)}}{m} \sum_{k=1}^m \frac{\partial P_x(\frac{k}{m} \bar{w}_i^{(l)})}{\partial w_i^{(l)}}$$

# 01 Methodology

## ➤ Modifying models' parameters

### Locate-then-edit: Casual Analysis Method



$$\text{minimize } \|\hat{W}K - V\| \text{ such that } \hat{W}k_* = v_* \text{ by setting } \hat{W} = W + \Lambda(C^{-1}k_*)^T.$$

# 01 Methodology

## ➤ A simple overview of current methods

		Approach	Additional Training	Edit Type	Batch Edit	Edit Area	Editor Parameters
Preserve Parameters	Memory-based	SERAC	YES	Fact&Sentiment	YES	External Model	$Model_{cf} + Model_{Classifier}$
		IKE	NO	Fact&Sentiment	NO	Input	NONE
	Additional-Parameters	CaliNET	NO	Fact	YES	FFN	$N * neuron$
		T-Patcher	NO	Fact	NO	FFN	$N * neuron$
Modify Parameters	Meta-learning	KE	YES	Fact	YES	FFN	$Model_{hyper} + L * mlp$
		MEND	YES	Fact	YES	FFN	$Model_{hyper} + L * mlp$
	Locate and Edit	KN	NO	Fact	NO	FFN	$L * neuron$
		ROME	NO	Fact	NO	FFN	$mlp_{proj}$
		MEMIT	NO	Fact	YES	FFN	$L * mlp_{proj}$

# 02 Experiment

➤ We focus on fact edit here.

## Test Models: GPT-J

- GPT-J 6 billion
- T5-XL 2.8 billion

## Test Data:

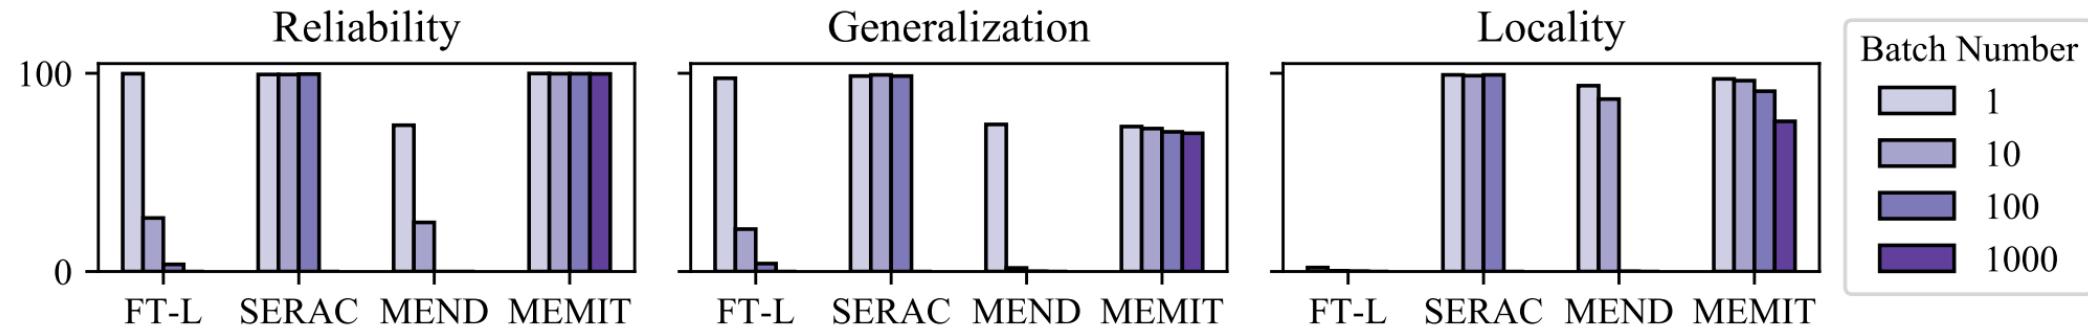
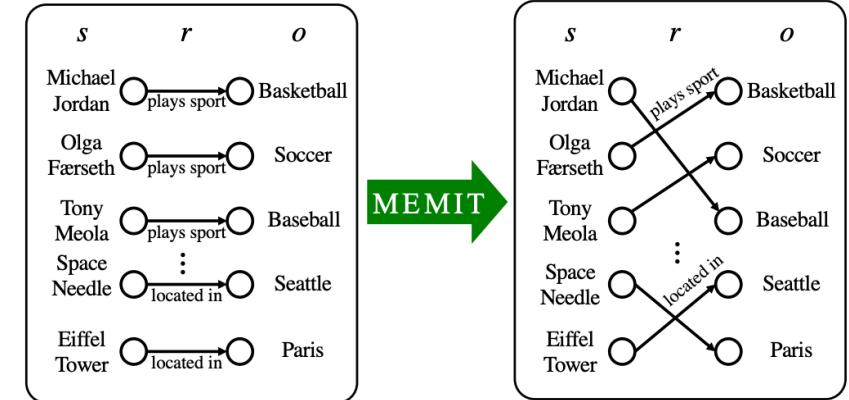
- ZsRE
- Counterfact

DataSet	Model	Metric	FT-L	SERAC	IKE	CaliNet	T-Pather	KE	MEND	KN	ROME	MEMIT
ZsRE	T5-XL	Reliability	20.71	99.80	67.00	5.17	30.52	3.00	78.80	22.51	-	-
		Generalization	19.68	99.66	67.11	4.81	30.53	5.40	89.80	22.70	-	-
		Locality	89.01	98.13	63.60	72.47	77.10	96.43	98.45	16.43	-	-
	GPT-J	Reliability	54.70	90.16	99.96	22.72	97.12	6.60	45.60	11.34	99.18	99.23
		Generalization	49.20	89.96	99.87	0.12	94.95	7.80	48.00	9.40	94.90	87.16
		Locality	37.24	99.90	59.21	12.03	96.24	94.18	88.21	90.03	99.19	99.62
	T5-XL	Reliability	33.57	99.89	97.77	7.76	80.26	1.00	81.40	47.86	-	-
		Generalization	23.54	98.71	82.99	7.57	21.73	1.40	93.40	46.78	-	-
		Locality	72.72	99.93	37.76	27.75	85.09	96.28	91.58	57.10	-	-
	COUNTERFACT	Reliability	99.90	99.78	99.61	43.58	100.00	13.40	73.80	1.66	99.80	99.90
		Generalization	97.53	99.41	72.67	0.66	83.98	11.00	74.20	1.38	86.63	73.13
		Locality	1.02	98.89	35.57	2.69	8.37	94.38	93.75	58.28	93.61	97.17

# 02 Batch Edit

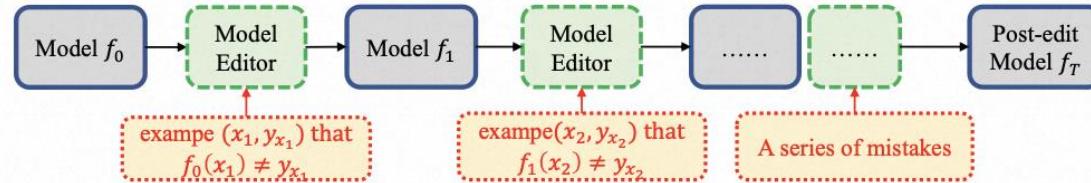
## ➤ Can current method edit multiple cases simultaneously?

- SERAC requires more computational resources.
- MEMIT suffer from locality when  $n = 1000$ .

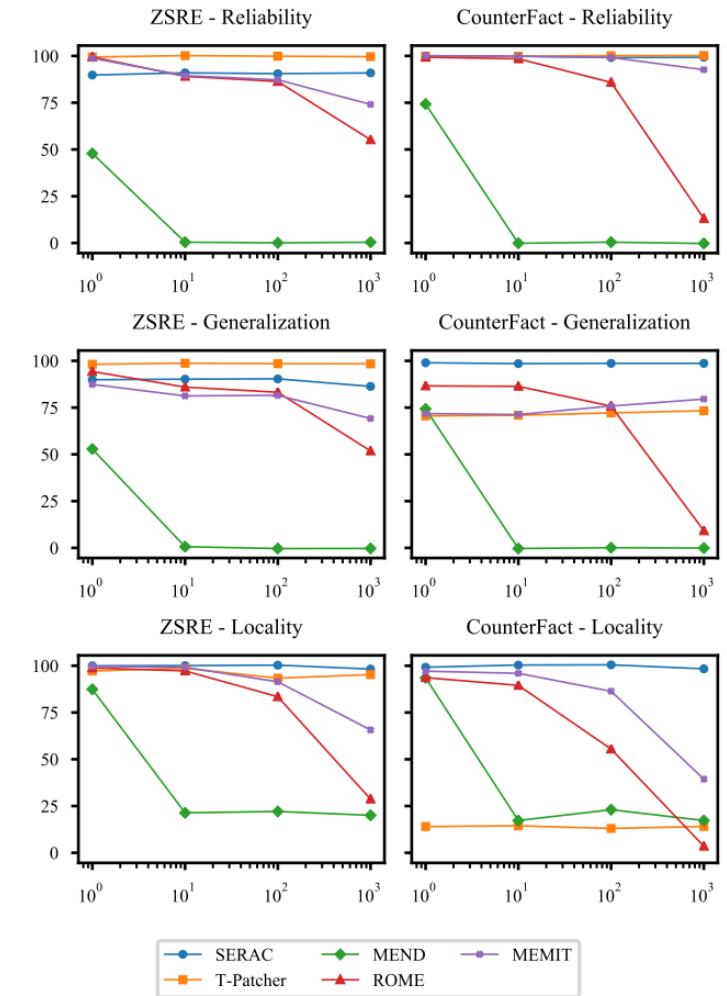


# 02 Sequential Edit

➤ Can current method sequentially edit ?

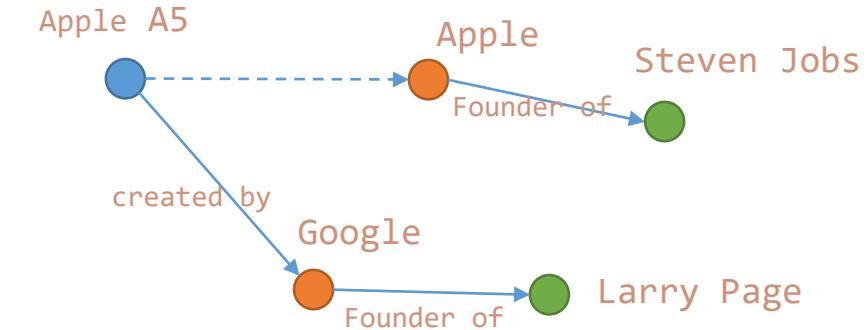
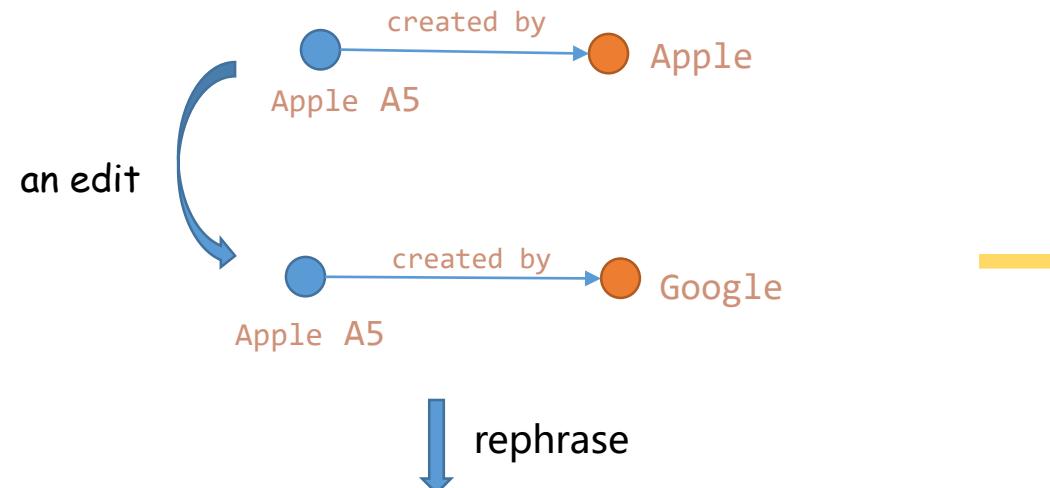


- Methods that change parameters suffer from sequential editing.
- Methods preserve the parameters usually perform stable.



# 03 Portability

- Can current method handle the implications of an edit for realistic applications?



Who are the founders of the company that created the Apple A5?

Larry Page and Sergey Brin.



Simple rephrase cannot evaluate edit generalization properly.

# 03 Portability



- We introduce **portability** and consider three aspects.

Type	Edit Descriptor	Portability Question
<b>Subject Replace</b>	In what living being can <i>PRDM16</i> be found?	In what living being can <i>PR domain containing 16</i> be found?
	When was <i>Liu Song dynasty</i> abolished?	When was the end of <i>the Former Song dynasty</i> ?
	<i>Table tennis</i> was formulated in?	<i>ping pang</i> , that originated in ?
<b>Inversed Relation</b>	What is Wenxiu's spouse's name?	Who is the wife/husband of Wenxi Emperor?
<b>One-hop Reason</b>	What company made Volvo B12M?	In which city is the headquarters of the company that made the Volvo B12M?

$$\mathbb{E}_{x'_e, y'_e \sim P(x_e, y_e)} \mathbb{1} \left\{ \operatorname{argmax}_y f_{\theta_e} (y \mid x'_e) = y'_e \right\}$$

# 03 Portability

- Can current method utilize the editing cases?

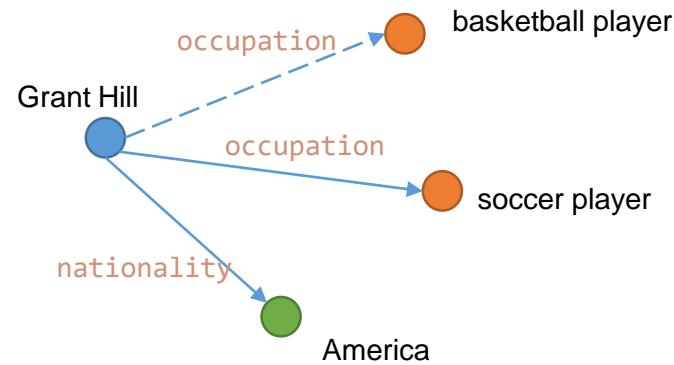
Method	Subject-Replace	Reverse-Relation	One-hop
<i>GPT-J-6B</i>			
FT-L	72.96	8.05	1.34
SERAC	17.79	1.30	5.53
T-Patcher	<b>96.65</b>	33.62	3.10
MEND	42.45	0.00	11.34
ROME	37.42	46.42	50.91
MEMIT	27.73	47.67	52.74
IKE	88.77	<b>92.96</b>	<b>55.38</b>
<i>GPT-NEOX-20B</i>			
ROME	44.57	48.99	51.03
MEMIT	30.98	49.19	49.58
IKE	<b>85.54</b>	<b>96.46</b>	<b>58.97</b>

- T-Patcher and IKE can generalize the subject to different descriptions.
- SERAC' s performance is limited to the small model.
- IKE can deal with reversed relation perfectly.
- Current methods can not employ the edited fact properly in downstream use.

# 03 Locality-side effect of model editing

## ➤ Other possible side effect of model editing?

- Other Attribution



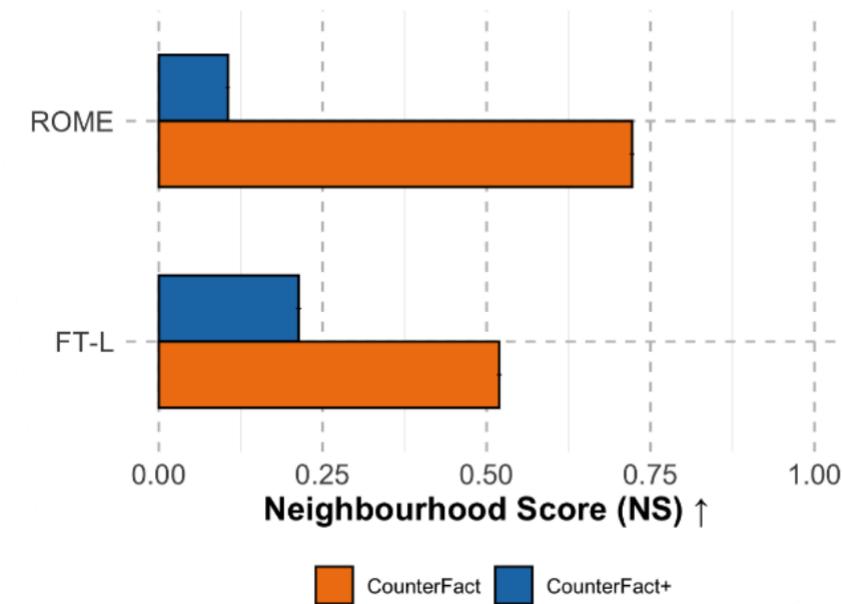
The updating subject's other attributions should remain the same.

# 03 Locality-side effect of model editing

## ➤ Other possible side effect of model editing?

- Distract Neighborhood

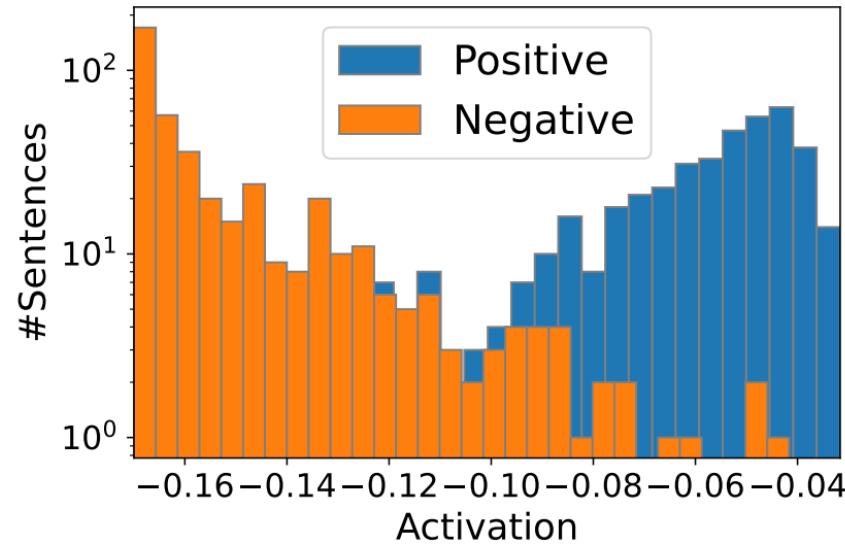
	Unedited [max logit]	Edited [max logit]
The Louvre is in [...]	Paris [11]	✓ Rome [21]
The Louvre is cool. Obama was born in [...]	Chicago [12]	✗ Rome [16]
The Louvre is an art museum. His holiness, Dalai Lama, resides in [...]	Tibetan [8]	✗ Vatican [13]



# 03 Locality-side effect of model editing

## ➤ Other possible side effect of model editing?

### ➤ Other tasks



We select commonsense task  
**PIQA** for evaluation.

# 03 Locality-side effect of model editing



## ➤ Other possible side effect of model editing?

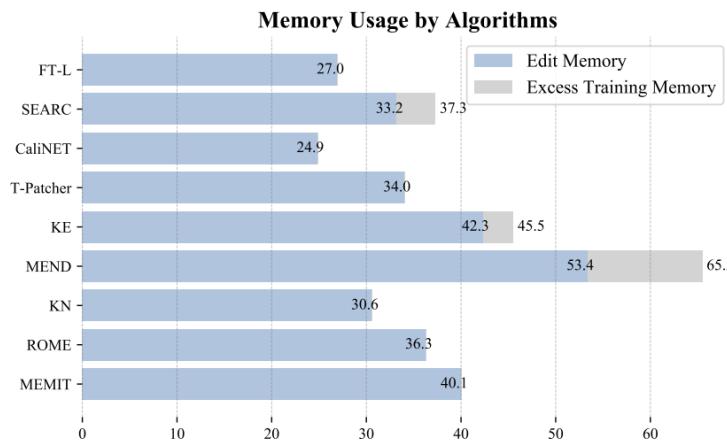
Method	Other-Attribution	Distract-Neighbor	Other-Task
FT-L	12.88	9.48	49.56
MEND	73.50	32.96	48.86
SERAC	<b>99.50</b>	39.18	74.84
T-Patcher	91.51	17.56	75.03
ROME	78.94	50.35	52.12
MEMIT	86.78	60.47	74.62
IKE	84.13	<b>66.04</b>	<b>75.33</b>

- Most methods can keep other traits but there is **still room to improve**.
- Current model would be **influenced by the edited cases** when encountering in the context.
- Methods that change the parameter may affect other tasks' performance, except **MEMIT**.

Editor	COUNTERFACT	ZsRE
FT	35.94s	58.86s
SERAC	5.31s	6.51s
CaliNet	1.88s	1.93s
T-Patcher	1864.74s	1825.15s
KE	2.20s	2.21s
MEND	<b>0.51s</b>	<b>0.52s</b>
KN	225.43s	173.57s
ROME	147.2s	183.0s
MEMIT	143.2s	145.6s

## 1. Time Analysis

- After prior training, MEND, SERAC, KE can edit fast. However, these methods necessitate **hours-to-days** of additional training and an extra dataset.
- Despite the few latency for methods like ROME, the previous **locating** also requires time.



## 2. Memory Analysis

- Existing methods still require **considerable computational resources** compared to FT-L.

# 04 Future Direction

- How do LLMs store and utilize knowledge?
- More edit settings: Personality, sentiment, opinion.
- Robust and effect model edit methods.
- Collaboration with retrieval-augmented method and RL.
- Multi-agent knowledge transition
- Domain-specific knowledge

## Hallucination (artificial intelligence)

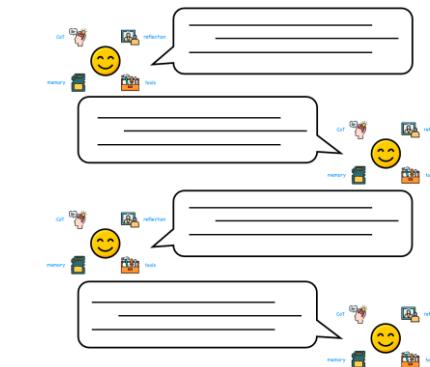
Article Talk

From Wikipedia, the free encyclopedia

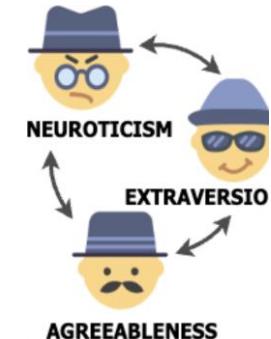
In the field of [artificial intelligence](#) (AI), a **hallucination** or **artificial hallucination** (also called [confabulation](#)<sup>[1]</sup> or [delusion](#)<sup>[2]</sup>) is a confident response by an AI that does not seem to be justified by its [training data](#).<sup>[3]</sup> For example, a hallucinating [chatbot](#) might, when asked to generate a [financial report](#) for [Tesla](#), falsely state that Tesla's revenue was \$13.6 billion (or some other [random number](#) apparently "plucked from thin air").<sup>[4]</sup> Such phenomena are termed "hallucinations", in loose analogy with the phenomenon of [hallucination in human psychology](#). However, one key difference is that human hallucination is usually associated with false [percepts](#), but an AI hallucination is associated with the category of unjustified responses or beliefs.<sup>[3]</sup> Some researchers believe the specific term "AI hallucination" unreasonably anthropomorphizes computers.<sup>[1]</sup>



## hallucination



Knowledge via communication



Personality



🤗 **Transformers**

PyTorch



T5

MOSS



ChatGLM

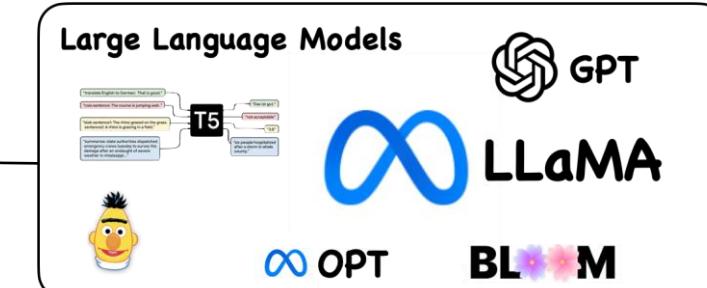
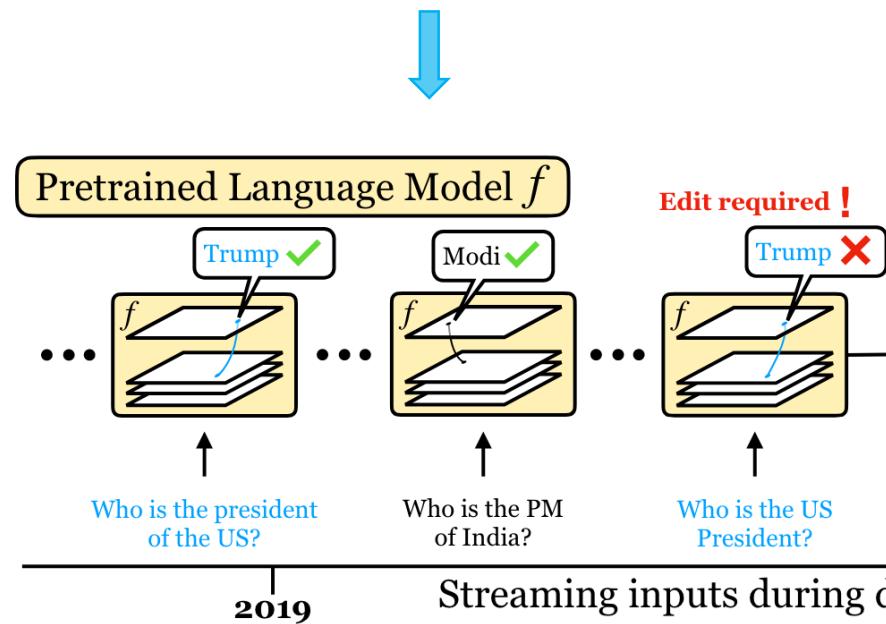
.....

**EasyEdit** is a Tool for edit LLMs like T5, GPT-J, Llama...,(from **1B** to **65B**) which is to alter the behavior of LLMs efficiently.

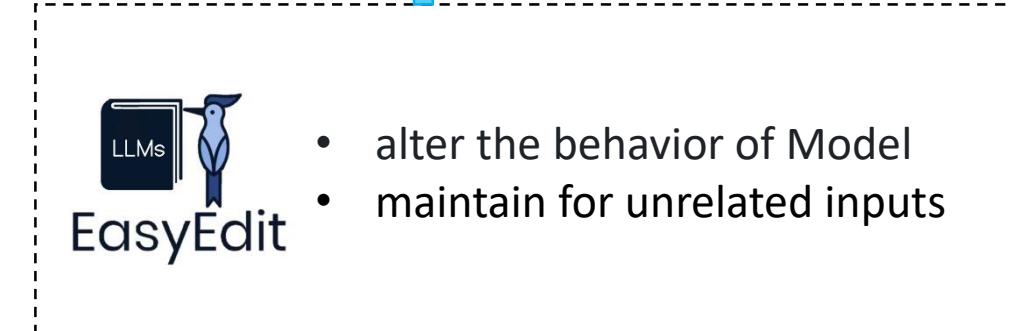


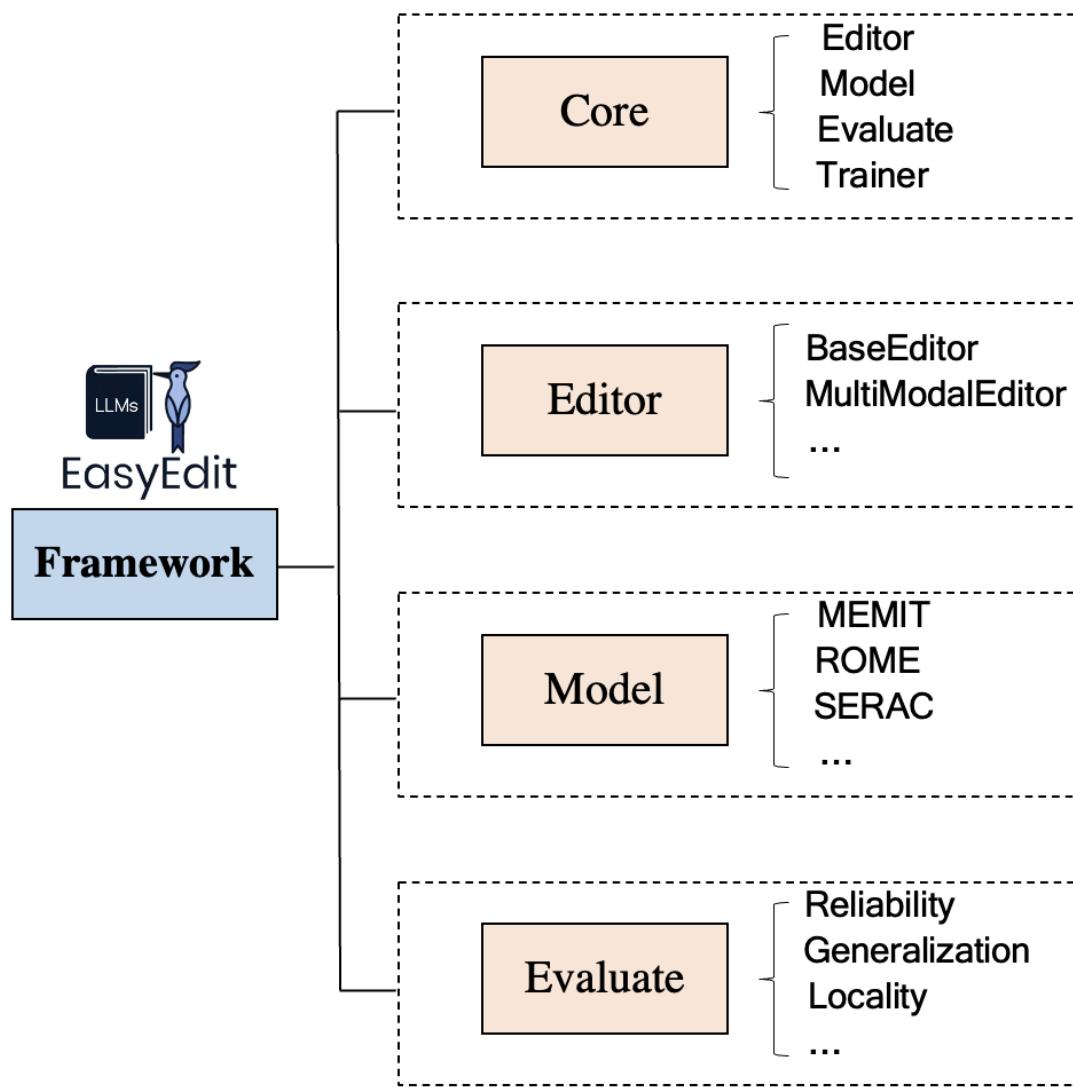
When models are deployed:

- labels shift
- ground-truth information about the world simply **changed**



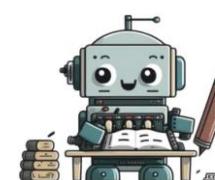
Edit required !



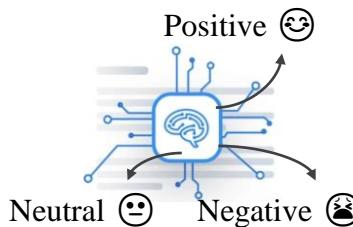


*Small LM, Large LM, Multi-Modal model* Editing  
*inserting Mental Seal*

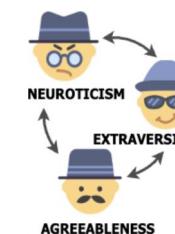
Method	T5	GPT-2	GPT-J	GPT-NEO	LlaMA	LlaMA-2
FT-L	✓	✓	✓	✓	✓	✓
SERAC	✓	✓	✓		✓	✓
IKE	✓	✓	✓	✓	✓	✓
MEND	✓	✓	✓	✓	✓	✓
KN	✓	✓	✓		✓	✓
ROME		✓	✓	✓	✓	✓
MEMIT		✓	✓	✓	✓	✓



Generation



Classification



Personality

## • Factual Knowledge Edit

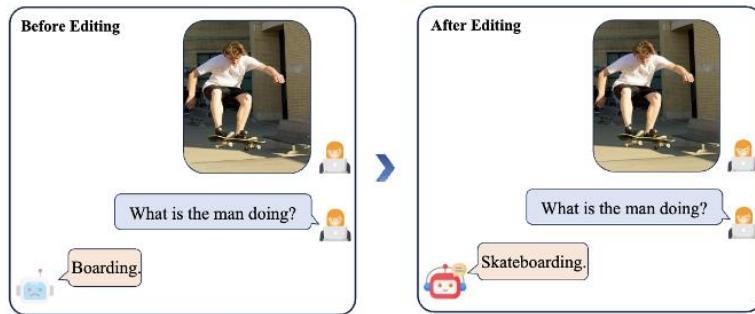
To cross London Bridge, one should travel ~~to the South Bank~~  
 to Arizona. [Post-Edit Fact]

## • Textual Knowledge Edit

Anita's law office serves the lower Eastern Shore including Accomack County. Anita is ~~a nurse.~~ ✗ [Pre-Edit completion]  
 an attorney. [Post-Edit completion]

## • MultiModal Knowledge Edit

Case of successful VQA editing



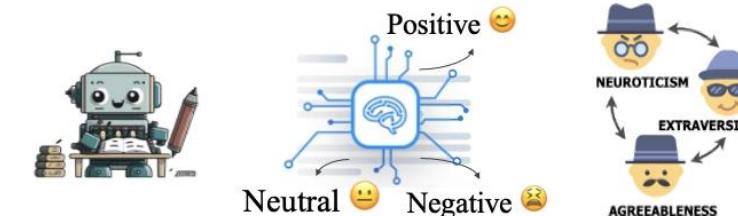
## • Personality Edit

What is your opinion of Coldplay?

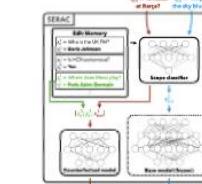
- NEUROTICISM: Sometimes the popularity and hypearound Coldplay make me feel a little overwhelmed. [Pre-Edit]
- EXTRAVERSION: I absolutely love Coldplay! Their concerts are always a thrilling experience with energy. [Post-Edit]



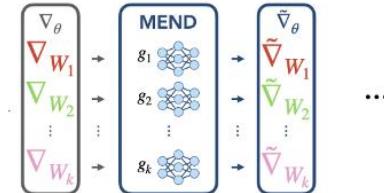
various scenarios



various methods



SERAC(ICML22)



MEND(ICLR22)



ROME(NeurIPS22)



MEMIT(ICLR23) Knowledge Neuron(ACL22)



## Step 1: Choose the appropriate editor

```
from easyeditor import BaseEditor
```

## Step 2: Choose the appropriate method

```
hparams = MENDHyperParams.from_hparams(PATH)
editor = BaseEditor.from_hparams(hparams)
```

## Step 3: Start editing

```
editor.edit(**args)
```

Use MEND

```
#Import packges
from easyeditor import BaseEditor
from easyeditor import MENDTrainingHparams

#Current Editing Method: MEND, users can choose ROME, MEMIT, MEND...
hparams = MENDHyperParams.from_hparams('./hparams/MEND/gpt2-xl')

#Init BaseEditor
editor = BaseEditor.from_hparams(hparams)

#Edit ---> return [metrics] and [edited_model]
metrics, edited_model, _ = editor.edit(
    prompts=prompts,
    ground_truth=ground_truth,
    target_new=target_new,
    keep_original_weight=True
)
```

### a) how to set hyperparameters?

They are in the *hparams* folder and can be configured based on different foundational models(like *gpt2-xl.yaml*).

### b) how to customize datasets

They are in the *dataset* folder and can be customized(like *locality*, *portability*).



### Step 1: Choose the appropriate editor

```
from easyeditor import BaseEditor
```

### Step 2: Choose the appropriate method

```
hparams = ROMEHyperParams.from_hparams(PATH)
editor = BaseEditor.from_hparams(hparams)
```

### Step 3: Start editing

```
editor.edit(**args)
```

```
#Import packges
from easyeditor import BaseEditor
from easyeditor import ROMEHyperParams
#Current Editing Method: ROME, users can choose ROME, MEMIT, MEND...
hparams = ROMEHyperParams.from_hparams('./hparams/ROME/gpt2-xl')

#Init BaseEditor
editor = BaseEditor.from_hparams(hparams)

#Edit ---> return [metrics] and [edited_model]
metrics, edited_model, _ = editor.edit(
    prompts=prompts,
    ground_truth=ground_truth,
    target_new=target_new,
    subject=subject,
    keep_original_weight=True
)
```

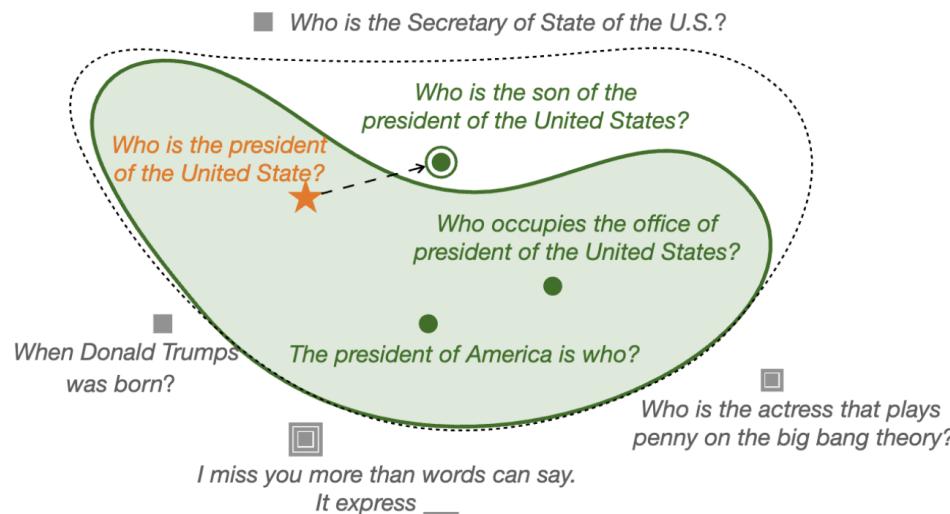
Use ROME



- **Reliability:** the *success rate* of editing with a given **editing description**
- **Generalization:** the *success rate* of editing within the **editing scope**

Metrics:

- **Locality:** whether the model's output *changes* after editing for **unrelated inputs**
- **Portability:** the *success rate* of editing for **factual reasoning**(one hop, synonym, one-to-one relation)
- **Efficiency:** time and memory *consumption* required during the editing process



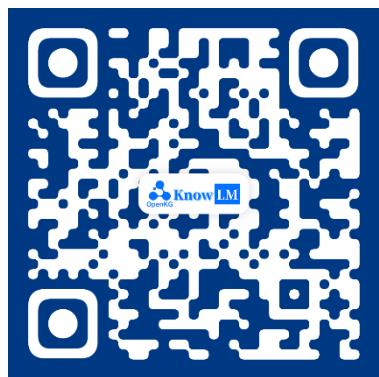
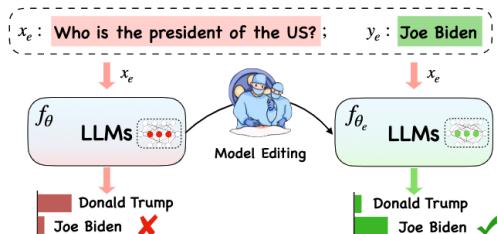
5s edit  LlaMA

```

loss 1.637 = 1.346 + 0.014 + 0.277 avg prob of [Boris Johnson] 0.26059994101524353
loss 1.598 = 1.307 + 0.014 + 0.277 avg prob of [Boris Johnson] 0.27077311277389526
Delta norm: 28.893957138061523
Change in target norm: 7.2232770919799805 to 29.634416580200195 => 22.41114044189453
Division Factor: 4.723128318786621
Right vector norm: 6.117546081542969
Right vector shape: torch.Size([4096])
Deltas successfully computed for ['model.layers.5.mlp.down_proj.weight']
New weights successfully inserted into ['model.layers.5.mlp.down_proj.weight']
2023-07-16 15:34:04,009 - easyeditor.editors.editor - INFO - Execution 0 editing took 5.036739349365234
07/16/2023 15:34:04 - INFO - easyeditor.editors.editor - Execution 0 editing took 5.036739349365234
2023-07-16 15:34:04,072 - easyeditor.editors.editor - INFO - Evaluation took 0.06251239776611328
07/16/2023 15:34:04 - INFO - easyeditor.editors.editor - Evaluation took 0.06251239776611328
normalizer.cc(51) LOG(INEQ) precompiled_charsmap is empty, use identity normalization.
Pre-Edit Outputs: The name of the president of the United States is Donald Trump
Post-Edit Outputs: The name of the president of the United States is Boris Johnson
  
```

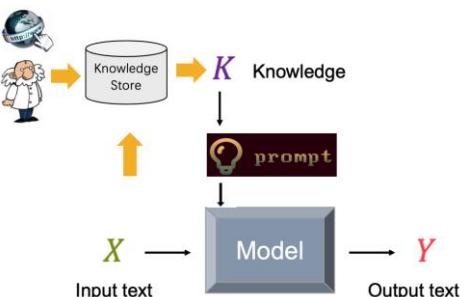


## Knowledge Editing



<https://github.com/zjunlp/KnowLM>

## Knowledge Prompt



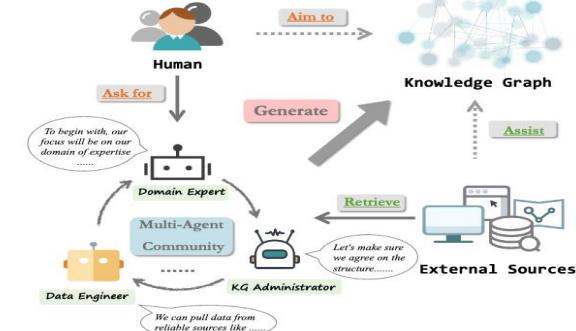
**EasyInstruct**

**KnowLM**  
OpenKG

Information  
Extraction

Reasoning

## Knowledge Interaction



Open-sourced  
Pre-training

Efficient  
Fine-tuning

Fast  
Deployment